

Rapidfire.py

```
from collections import Counter
```

```
from transformers import pipeline
```

```
from groq import Groq
```

```
from collections import Counter
```

```
from nltk.tokenize import word_tokenize
```

```
from nltk.corpus import stopwords
```

```
import json
```

```
import numpy as np
```

```
from .audio_process import process_audio_upload, convert_wav_to_mp3,  
transcribe_audio, save_uploaded_file
```

```
def generate_incomplete_analogy():
```

```
    client =
```

```
    Groq(api_key="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")
```

```
    completion = client.chat.completions.create(
```

```
        model="llama-3.3-70b-versatile",
```

```
        messages=[
```

```
            {
```

```
                "role": "system",
```

```
                "content": (
```

```
                    "You are a helpful AI Assistant. "
```

```
                    "Generate a single incomplete analogy prompt. "
```

```
                    "For example: 'Learning is like', 'Love is like'. "
```

```
                    "Output exactly one incomplete analogy without any additional words."
```

```
                )
```

```
            },
```

```
            {
```

```
                "role": "user",
```

```

        "content": "give a random incomplete analogy",
    }
],
temperature=2,
frequency_penalty=0.0,
max_completion_tokens=1024,
top_p=1,
stream=True,
stop=None,
)
print(completion)
# Handling the streamed output
response_content = ""
for chunk in completion:
    response_content = response_content+(chunk.choices[0].delta.content or "")
    print(chunk.choices[0].delta.content or "", end="")

return response_content

```

```

def extract_topic(transcript):

```

```

    # Initialize english stopwords
    english_stopwords = stopwords.words("english")

    #convert article to tokens
    tokens = word_tokenize(transcript)

    #extract alpha words and convert to lowercase

```

```
alpha_lower_tokens = [word.lower() for word in tokens if word.isalpha()]
```

```
#remove stopwords
```

```
alpha_no_stopwords = [word for word in alpha_lower_tokens if word not in  
english_stopwords]
```

```
#Count word
```

```
BoW = Counter(alpha_no_stopwords)
```

```
#Most common words
```

```
return BoW.most_common(3)
```

```
def score_analogy_with_groq(transcript, incomplete_analogy):
```

```
    """
```

```
    Use Groq to evaluate analogy relevance and creativity based on the transcript  
    and the incomplete analogy prompt.
```

```
Parameters:
```

```
    transcript (str): The transcribed text from the user's response
```

```
    incomplete_analogy (str): The incomplete analogy prompt (e.g., "Success is like")
```

```
Returns:
```

```
    dict: A dictionary with analogy_relevance and creativity scores
```

```
    """
```

```
    client =
```

```
Groq(api_key="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")
```

```
    prompt = f"""
```

```
    Below is an incomplete analogy prompt and a user's spoken response to complete it.
```

Incomplete Analogy: "{incomplete_analogy}"

User's Response: "{transcript}"

Please evaluate the response based on two criteria:

1. Relevance (0-10): How well does the response connect to the analogy prompt?

Does it create a clear and appropriate comparison?

2. Creativity (0-10): How original, insightful, or thought-provoking is the analogy?

Does it provide a fresh perspective or use unexpected connections?

Return your evaluation as a JSON object with two properties:

- analogy_relevance: a number between 0 and 10 (with up to 2 decimal places)

- creativity: a number between 0 and 10 (with up to 2 decimal places)

Response must be in this exact JSON format and nothing else:

```
{{  
  "analogy_relevance": 0.00,  
  "creativity": 0.00  
}}
```

"""

```
completion = client.chat.completions.create(  
  model="llama-3.3-70b-versatile",  
  messages=[  
    {  
      "role": "system",
```

```
      "role": "system",
```

```
        "content": "You are an AI assistant that evaluates analogies based on relevance and creativity. Provide numeric scores only."
```

```
    },  
    {  
        "role": "user",  
        "content": prompt  
    }  
],  
temperature=0.8,  
max_completion_tokens=256  
)
```

```
response_content = completion.choices[0].message.content
```

```
response_dict = json.loads(response_content)
```

```
return {  
    "analogy_relevance": round(response_dict['analogy_relevance'], 2),  
    "creativity": round(response_dict["creativity"], 2)  
}
```

```
def process_rapidfire_audio(audio_file_path, analogy):
```

```
    result = process_audio_upload(audio_file_path)
```

```
    segments = result.get("segments", [])
```

```
    transcript = result.get("text", "")
```

```
    # Calculate word-level timestamps
```

```
    word_timestamps = []
```

```
previous_word_end = None  
gaps = [] # store gaps between consecutive words
```

```
for segment in segments:
```

```
    seg_start = segment["start"]  
    seg_end = segment["end"]  
    seg_text = segment["text"].strip()  
    words = seg_text.split()  
    num_words = len(words)  
    if num_words == 0:  
        continue  
    duration = seg_end - seg_start  
    word_duration = duration / num_words
```

```
    for i, word in enumerate(words):
```

```
        word_start = seg_start + i * word_duration  
        word_end = word_start + word_duration  
        word_timestamps.append({  
            "word": word,  
            "start": round(word_start, 2),  
            "end": round(word_end, 2)  
        })
```

```
    if previous_word_end is not None:
```

```
        gap = word_start - previous_word_end  
        gaps.append(gap)  
        previous_word_end = word_end
```

```
if transcript.strip() == "":
```

```

    speech_continuity = 0
else:
    if gaps:
        gaps_array = np.array(gaps)
        avg_gap = np.mean(gaps_array)
        std_gap = np.std(gaps_array)

        # Penalize both large average gaps and inconsistent timing
        speech_continuity = max(0, 10 - (avg_gap * 5) - (std_gap * 3))
    else:
        speech_continuity = 10.0

# Score analogy relevance and creativity using Groq
analogy_scores = score_analogy_with_groq(transcript, analogy)
analogy_relevance = analogy_scores["analogy_relevance"]
creativity = analogy_scores["creativity"]

# Extract text topic using NLP
text_topic = extract_topic(transcript)

# Calculate overall score
overall_rapidfire_score = (speech_continuity + analogy_relevance + creativity) / 3

metrics = {
    "speech_continuity": round(speech_continuity, 2),
    "analogy_relevance": analogy_relevance,
    "creativity": creativity,
    "overall_rapidfire_score": round(overall_rapidfire_score, 2),

```

```
    "text_topic": text_topic
}

return {
    "transcript": transcript,
    "word_timestamps": word_timestamps,
    "metrics": metrics,
    "generated_analogy": analogy
}
```


Triplestep.py

```
import os
```

```
from pydub import AudioSegment
```

```
import whisper
```

```
import random
```

```
from groq import Groq
```

```
import nltk
```

```
from sentence_transformers import SentenceTransformer, util
```

```
from nltk.tokenize import sent_tokenize
```

```
import numpy as np
```

```
import nltk
```

```
from nltk.tokenize import word_tokenize
```

```
import nltk
```

```
from groq import Groq
```

```
def analyze_distractor_smoothness(transcript: str, distractor_words: list) -> float:
```

```
    """
```

```
    Analyze how smoothly distractor words are integrated in the transcript.
```

```
    For each sentence containing any distractor word (case-insensitive), compute the cosine similarity
```

```
    between that sentence and its adjacent sentences (previous and next). A higher similarity indicates smoother integration.
```

```
    Returns a smoothness score between 0 and 10. If no sentence contains a distractor word, returns 10.0.
```

```
    """
```

```
    sentences = sent_tokenize(transcript)
```

```
    if not sentences:
```

```
return 0.0
```

```
model = SentenceTransformer('all-mpnet-base-v2')
```

```
distracted_similarities = []
```

```
for i, sentence in enumerate(sentences):
```

```
    if any(dw.lower() in sentence.lower() for dw in distractor_words):
```

```
        neighbor_sims = []
```

```
        # Compute similarity with previous sentence if exists
```

```
        if i > 0:
```

```
            try:
```

```
                sim_prev = util.cos_sim(
```

```
                    model.encode(sentence, convert_to_tensor=True),
```

```
                    model.encode(sentences[i-1].strip(), convert_to_tensor=True)
```

```
                ).item()
```

```
                neighbor_sims.append(sim_prev)
```

```
            except Exception as e:
```

```
                print("Error computing similarity with previous sentence:", e)
```

```
        # Compute similarity with next sentence if exists
```

```
        if i < len(sentences) - 1:
```

```
            try:
```

```
                sim_next = util.cos_sim(
```

```
                    model.encode(sentence, convert_to_tensor=True),
```

```
                    model.encode(sentences[i+1].strip(), convert_to_tensor=True)
```

```
                ).item()
```

```
                neighbor_sims.append(sim_next)
```

```
            except Exception as e:
```

```
                print("Error computing similarity with next sentence:", e)
```

```
if neighbor_sims:
    distracted_similarities.append(np.mean(neighbor_sims))
```

```
if not distracted_similarities:
    return 10.0
```

```
avg_similarity = np.mean(distracted_similarities)
smoothness_score = avg_similarity * 10
smoothness_score = max(0, min(smoothness_score, 10))
return round(smoothness_score, 2)
```

```
def analyze_topic_adherence(transcript: str, expected_topic: str) -> float:
```

```
    """
```

Analyze topic adherence by:

1. Splitting the transcript into sentences.
2. Computing embeddings for each non-empty sentence using SentenceTransformer.
3. Computing the embedding for the expected topic.
4. Calculating the cosine similarity between each sentence and the expected topic.
5. Aggregating these similarities into a global topic adherence score (0-10).

Returns a score between 0 and 10.

```
    """
```

```
sentences = sent_tokenize(transcript)
```

```
if not sentences:
```

```
    return 0.0
```

```
model = SentenceTransformer('all-mpnet-base-v2')
```

```
try:
```

```

    topic_embedding = model.encode(expected_topic, convert_to_tensor=True)
except Exception as e:
    print("Error encoding expected_topic:", e)
    return 0.0

similarities = []
for sentence in sentences:
    sentence = sentence.strip()
    if not sentence:
        continue
    try:
        sentence_embedding = model.encode(sentence, convert_to_tensor=True)
    except Exception as e:
        print("Error encoding sentence:", sentence, e)
        continue
    # Check if embedding is valid (non-empty)
    if sentence_embedding is None or sentence_embedding.shape[0] == 0:
        print("Empty embedding for sentence:", sentence)
        continue
    try:
        cos_sim = util.cos_sim(topic_embedding, sentence_embedding).item()
    except Exception as e:
        print("Error computing cosine similarity for sentence:", sentence, e)
        continue
    similarities.append(cos_sim)

if not similarities:
    return 0.0

```

```
avg_similarity = np.mean(similarities)
```

```
score = avg_similarity * 10
```

```
score = max(0, min(score, 10))
```

```
return round(score, 2)
```

```
def analyze_coherence(transcript):
```

```
    """
```

Analyze speech coherence by:

1. Segmenting the transcript into sentences.
2. Embedding each sentence using Sentence-BERT.
3. Computing cosine similarity between adjacent sentence embeddings.
4. Aggregating these similarities into a global coherence score.

Returns:

- coherence_score: A score (0-10) representing global coherence.
- sentences: List of segmented sentences.
- similarities: List of cosine similarities between adjacent sentences.

```
    """
```

```
# Sentence Segmentation
```

```
sentences = nltk.sent_tokenize(transcript)
```

```
if len(sentences) < 2:
```

```
    return 10.0, sentences, []
```

```
# Embed Sentences with Sentence-BERT
```

```
model = SentenceTransformer('all-mpnet-base-v2')
```

```
sentence_embeddings = model.encode(sentences, convert_to_tensor=True)
```

```
# Compute Cosine Similarity Between Adjacent Sentences
```

```
similarities = []
```

```
for i in range(len(sentences) - 1):
```

```
    cos_sim = util.cos_sim(sentence_embeddings[i], sentence_embeddings[i+1]).item()
```

```
    similarities.append(cos_sim)
```

```
# Aggregate Similarities: Compute the average similarity
```

```
avg_similarity = np.mean(similarities)
```

```
coherence_score = avg_similarity * 10
```

```
return coherence_score
```

```
def save_uploaded_file(audio_file_path, destination_path):
```

```
    with open(audio_file_path, 'rb') as in_file:
```

```
        data = in_file.read()
```

```
    with open(destination_path, 'wb') as out_file:
```

```
        out_file.write(data)
```

```
def convert_wav_to_mp3(wav_file, mp3_file="output.mp3"):
```

```
    try:
```

```
        audio = AudioSegment.from_wav(wav_file)
```

```
        audio.export(mp3_file, format="mp3")
```

```
        print(f"Converted {wav_file} to {mp3_file}")
```

```
        return mp3_file
```

```
    except Exception as e:
```

```
        print("Error converting WAV to MP3:", e)
```

```
        raise
```

```
def transcribe_audio(mp3_file):  
    print("Loading Whisper model...")  
    model = whisper.load_model("small.en")  
    print("Transcribing audio...")  
    result = model.transcribe(mp3_file)  
    return result["text"]
```

```
def process_audio_upload(audio_file_path):  
    wav_path = "temp_output.wav"  
    mp3_path = "temp_output.mp3"
```

```
    try:  
        print("Converting webm to wav...")  
        audio_segment = AudioSegment.from_file(audio_file_path, format="webm")  
        audio_segment.export(wav_path, format="wav")  
    except Exception as e:  
        print("Error converting webm to wav:", e)  
        raise
```

```
    convert_wav_to_mp3(wav_path, mp3_path)
```

```
    transcript = transcribe_audio(mp3_path)
```

```
    for file in [wav_path, mp3_path]:  
        if os.path.exists(file):  
            os.remove(file)
```

```
return transcript
```

```
def generate_topics():
```

```
    """
```

```
    Generate a main speaking topic and contextually relevant distractor words.
```

```
    The response from the model is expected to contain lines like:
```

```
        Main Topic: <topic text>
```

```
        Distractor Words: <word1>, <word2>, <word3>
```

```
    Returns:
```

```
        dict: {"main_topic": <str>, "distractor_words": [<str>, ...]}
```

```
    """
```

```
    client =
```

```
    Groq(api_key="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")
```

```
    completion = client.chat.completions.create(
```

```
        model="llama-3.3-70b-versatile",
```

```
        messages=[
```

```
            {
```

```
                "role": "system",
```

```
                "content": ("You are a helpful AI Assistant. You should generate and display a main  
speaking topic, like Cars tourism travelling "
```

```
                    "and then generate contextually relevant distractor words related to that  
topic. "
```

```
                    "Please output the result in the following format exactly:\n\n"
```

```
                    "Main Topic: <your topic here>\n"
```

```
                    "Distractor Words: <word1>, <word2>, <word3>")
```

```
            },
```



```

    {
        "role": "user",
        "content": "Give a main topic and some distractor words."
    }
],
temperature=2,
frequency_penalty=0.0,
max_completion_tokens=1024,
top_p=1,
stream=True,
stop=None,
)

```

```

response_content = ""
for chunk in completion:
    # Each chunk's delta content may be None; we append if present.
    delta_text = chunk.choices[0].delta.content or ""
    response_content += delta_text
    print(delta_text, end="")

```

```

main_topic = ""
distractor_words = []
lines = response_content.splitlines()
for line in lines:
    if line.lower().startswith("main topic:"):
        main_topic = line.split(":", 1)[1].strip()
    elif line.lower().startswith("distractor words:"):
        words_str = line.split(":", 1)[1].strip()

```

```

    distractor_words = [w.strip() for w in words_str.split(",") if w.strip()]

    return {"main_topic": main_topic, "distractor_words": distractor_words}

def process_triple_step_audio(audio_file_path, main_topic, distractor_words):

    transcript = process_audio_upload(audio_file_path)

    # evaluation scores
    coherence_result = analyze_coherence(transcript)
    if isinstance(coherence_result, tuple):
        coherence_score = coherence_result[0]
    else:
        coherence_score = coherence_result

    topic_adherence_score = analyze_topic_adherence(transcript, main_topic)
    distraction_handling_score = analyze_distractor_smoothness(transcript, distractor_words)

    overall_triple_step_score = (coherence_score + topic_adherence_score +
    distraction_handling_score) / 3.0

    result = {
        "main_topic": main_topic,
        "distractor_words": distractor_words,
        "transcript": transcript,
        "coherence_score": coherence_score,
        "topic_adherence_score": topic_adherence_score,
        "distraction_handling_score": distraction_handling_score,

```

```
    "overall_triple_step_score": overall_triple_step_score
}

return result
```

Conductor.py

import os

from pydub import AudioSegment

import whisper

from groq import Groq

import librosa

import numpy as np

import nltk

import joblib

def save_uploaded_file(audio_file_path, destination_path):

with open(audio_file_path, 'rb') as in_file:

data = in_file.read()

with open(destination_path, 'wb') as out_file:

out_file.write(data)

def convert_wav_to_mp3(wav_file, mp3_file="output.mp3"):

try:

audio = AudioSegment.from_wav(wav_file)

audio.export(mp3_file, format="mp3")

print(f"Converted {wav_file} to {mp3_file}")

return mp3_file

except Exception as e:

print("Error converting WAV to MP3:", e)

raise

def transcribe_audio(mp3_file):

print("Loading Whisper model...")

```

model = whisper.load_model("small.en")
print("Transcribing audio...")
result = model.transcribe(mp3_file)
return result["text"]

```

```

def process_audio_upload(audio_file_path):
    wav_path = "temp_output.wav"
    mp3_path = "temp_output.mp3"
    try:
        print("Converting webm to wav...")
        audio_segment = AudioSegment.from_file(audio_file_path, format="webm")
        audio_segment.export(wav_path, format="wav")
    except Exception as e:
        print("Error converting webm to wav:", e)
        raise

    convert_wav_to_mp3(wav_path, mp3_path)
    transcript = transcribe_audio(mp3_path)
    for file in [wav_path, mp3_path]:
        if os.path.exists(file):
            os.remove(file)
    return transcript

```

```

def generate_conductor_exercise():

```

```

    """

```

Generate instructions for an exercise to improve vocal variety and expression.

The response from the model is expected to contain lines like:

Energy Levels: <energy_level1>, <energy_level2>, <energy_level3>

Moods: <mood1>, <mood2>, <mood3>

Improvement Suggestions: <suggestion1>, <suggestion2>, <suggestion3>

Returns:

```
dict: {  
    "energy_levels": [<str>, ...],  
    "moods": [<str>, ...],  
    "improvement_suggestions": [<str>, ...]  
}
```

"""

client =

Groq(api_key="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")

completion = client.chat.completions.create(

model="llama-3.3-70b-versatile",

messages=[

{

"role": "system",

"content": (

"You are a helpful AI Assistant. Generate an exercise prompt for improving vocal variety and expression. "

"The exercise should guide users through different energy levels and moods, and include instructions for real-time voice analysis "

"to track energy levels, analyze vocal variety, provide instant feedback on mood matching, and generate personalized improvement suggestions. "

"Please output the result in the following format exactly:\n\n"

"Energy level should only be High, Medium, or Low.\n"

```
"Mood should only be Joy, Sadness, Fear, Anger, Surprise, Neutral, Disgust, or Shame.\n"
```

```
"Energy Levels: <energy_level1>, <energy_level2>, <energy_level3>\n"
```

```
"Moods: <mood1>, <mood2>, <mood3>\n"
```

```
"Improvement Suggestions: <suggestion1>, <suggestion2>, <suggestion3>\n\n"
```

```
"Note: Moods must be chosen only from the following values: joy, sadness, fear, anger, surprise, neutral, disgust, shame."
```

```
)  
  
,  
  
{  
  
    "role": "user",  
  
    "content": "Generate an exercise prompt for improving vocal variety and expression."  
  
}  
  
],  
  
temperature=0.5,  
frequency_penalty=0.0,  
max_completion_tokens=1024,  
top_p=1,  
stream=True,  
stop=None,  
)
```

```
response_content = ""
```

```
for chunk in completion:
```

```
    # Append each delta's content (if present) to the response_content.
```

```
    delta_text = chunk.choices[0].delta.content or ""
```

```
    response_content += delta_text
```

```
print(delta_text, end="") # print for debugging
```

```
energy_levels = []
```

```
moods = []
```

```
improvement_suggestions = []
```

```
lines = response_content.splitlines()
```

```
for line in lines:
```

```
    if line.lower().startswith("energy levels:"):

```

```
        levels_str = line.split(":", 1)[1].strip()

```

```
        energy_levels = [lvl.strip() for lvl in levels_str.split(",") if lvl.strip()]

```

```
    elif line.lower().startswith("moods:"):

```

```
        moods_str = line.split(":", 1)[1].strip()

```

```
        moods = [m.strip() for m in moods_str.split(",") if m.strip()]

```

```
    elif line.lower().startswith("improvement suggestions:"):

```

```
        suggestions_str = line.split(":", 1)[1].strip()

```

```
        improvement_suggestions = [s.strip() for s in suggestions_str.split(",") if s.strip()]

```

```
return {

```

```
    "energy_levels": energy_levels,

```

```
    "moods": moods,

```

```
    "improvement_suggestions": improvement_suggestions

```

```
}
```

```
def remove_adjacent_duplicates(seq):

```

```
    if not seq:

```

```
        return []

```

```
    result = [seq[0]]

```



```
for item in seq[1:]:
    if item != result[-1]:
        result.append(item)
return result
```

```
def score_sequence_match(audio_sequence, target_sequence):
    # Remove adjacent duplicates
    audio_sequence = remove_adjacent_duplicates(audio_sequence)

    if not target_sequence or not audio_sequence:
        return 0.0

    # Convert categorical values to numerical
    energy_map = {"low": 0, "medium": 1, "high": 2}

    # Convert sequences to numerical values
    num_audio = [energy_map.get(level.lower(), 1) for level in audio_sequence]
    num_target = [energy_map.get(level.lower(), 1) for level in target_sequence]

    # This allows for partial matching and timing flexibility
    max_score = len(target_sequence)
    score = 0.0

    i, j = 0, 0
    while i < len(num_audio) and j < len(num_target):
        # Exact match
        if num_audio[i] == num_target[j]:
            score += 1.0
```

```

# Close match (off by one level)
elif abs(num_audio[i] - num_target[j]) == 1:
    score += 0.5

# Move forward in sequences
if i < len(num_audio) - 1 and j < len(num_target) - 1:
    # Determine which sequence to advance
    if num_audio[i+1] == num_target[j]:
        i += 1
    elif num_audio[i] == num_target[j+1]:
        j += 1
    else:
        i += 1
        j += 1
else:
    i += 1
    j += 1

# Normalize to 0-10 scale
final_score = (score / max_score) * 10

return max(round(final_score, 2), 10)

def analyze_mood_matches(transcript, target_moods_list):

    allowed_moods = ["joy", "sadness", "fear", "anger", "surprise", "neutral", "disgust",
"shame"]

    sentences = nltk.sent_tokenize(transcript)

```

```
results = []
```

```
current_dir = os.path.dirname(os.path.abspath(__file__))
```

```
pipeline_path = os.path.join(current_dir, "emotion_classifier_pipe_lr.pkl")
```

```
# Verify file exists
```

```
if not os.path.exists(pipeline_path):
```

```
    raise FileNotFoundError(f"Classifier pipeline not found at: {pipeline_path}")
```

```
# Load the classifier
```

```
with open(pipeline_path, "rb") as pipeline_file:
```

```
    loaded_pipe_lr = joblib.load(pipeline_file)
```

```
# Related mood groups for partial matching
```

```
related_moods = {
```

```
    "joy": ["surprise"],
```

```
    "sadness": ["shame", "disgust"],
```

```
    "fear": ["surprise", "shame"],
```

```
    "anger": ["disgust"],
```

```
    "surprise": ["joy", "fear"],
```

```
    "neutral": [],
```

```
    "disgust": ["anger", "sadness"],
```

```
    "shame": ["sadness", "fear"]
```

```
}
```

```
total_score = 0.3
```

```
total_sentences = 1
```

```

for i, sentence in enumerate(sentences):

    sentence = sentence.strip()

    if not sentence:

        continue

    total_sentences += 1

    # Get emotion prediction for this sentence
    predicted_emotion = loaded_pipe_lr.predict([sentence])[0].lower()

    # If we have a target mood for this sentence
    if target_moods_list and i < len(target_moods_list):

        expected_mood = target_moods_list.lower()

        # Exact match
        if predicted_emotion in expected_mood:

            total_score += 1.0

        # Related mood (partial match)
        elif predicted_emotion in related_moods.get(expected_mood, []):

            total_score += 0.5

        # Check if expected mood is in the related moods of predicted emotion
        elif expected_mood in related_moods.get(predicted_emotion, []):

            total_score += 0.3

    results.append({

        "sentence": sentence,

        "predicted_mood": predicted_emotion

    })

```

```
# Normalize score to 0-10 scale
print("MOOD sequence", results)
mood_score = (total_score/ total_sentences) * 10
return max(round(mood_score, 2),10)
```

```
def get_energy_level_sequence(audio_path, sr=22050, segment_duration=1.0):
    # Load audio
    y, sr = librosa.load(audio_path, sr=sr)

    # Compute RMS energy over frames
    rms = librosa.feature.rms(y=y)[0]
    hop_length = 512 # default hop_length in librosa.feature.rms
    times = librosa.frames_to_time(np.arange(len(rms)), sr=sr, hop_length=hop_length)

    # Divide audio into segments of 'segment_duration' seconds.
    max_time = times[-1]
    num_segments = int(np.ceil(max_time / segment_duration))
    segment_energies = []

    for i in range(num_segments):
        start_time = i * segment_duration
        end_time = (i + 1) * segment_duration
        indices = np.where((times >= start_time) & (times < end_time))[0]
        if len(indices) == 0:
            avg_energy = 0.0
        else:
            avg_energy = np.mean(rms[indices])
```

```

    segment_energies.append(avg_energy)

segment_energies = np.array(segment_energies)

# Define thresholds using quantiles
low_threshold = np.quantile(segment_energies, 0.33)
high_threshold = np.quantile(segment_energies, 0.66)

# Map each segment's average energy to a category.
energy_sequence = []
for energy in segment_energies:
    if energy < low_threshold:
        energy_sequence.append("low")
    elif energy < high_threshold:
        energy_sequence.append("medium")
    else:
        energy_sequence.append("high")
return energy_sequence

def generate_targeted_suggestions(energy_score, mood_score, audio_sequence,
target_sequence):
    """
    Generate targeted improvement suggestions based on actual performance
    """
    suggestions = []

    # Energy-related suggestions
    if energy_score < 5.0:

```

```

if len(set(audio_sequence)) <= 1:
    suggestions.append("Try using more varied energy levels - your delivery was mostly at
one level")
else:
    # Check if specific energy levels are missing
    audio_levels = set(level.lower() for level in audio_sequence)
    target_levels = set(level.lower() for level in target_sequence)
    missing_levels = target_levels - audio_levels

    if "high" in missing_levels:
        suggestions.append("Work on incorporating higher energy moments in your
delivery")
    if "low" in missing_levels:
        suggestions.append("Practice including quieter, more intimate moments in your
delivery")

# Mood-related suggestions
if mood_score < 5.0:
    suggestions.append("Focus on matching your vocal tone to the intended emotion")
    suggestions.append("Try exaggerating the emotional quality to make it more
recognizable")

# General suggestions if we don't have enough targeted ones
if len(suggestions) < 2:
    general_suggestions = [
        "Record yourself and listen back to identify subtle mood inconsistencies",
        "Try mirroring professional speakers to develop better vocal variety",
        "Work on maintaining consistent volume while varying your pitch and pace"
    ]

```

```

        # Add general suggestions until we have at least 2
        while len(suggestions) < 2 and general_suggestions:
            suggestions.append(general_suggestions.pop(0))

    return suggestions[:2] # Return top 2 suggestions

def process_conductor_audio(audio_file_path, instructions, energy_levels, moods):
    """
    Process the Conductor exercise audio with improved scoring logic
    """
    transcript = process_audio_upload(audio_file_path)

    # Get energy sequence from audio
    audio_sequence = get_energy_level_sequence(audio_file_path, segment_duration=1.0)

    # Calculate energy level score
    energy_level_score = score_sequence_match(audio_sequence, energy_levels)

    # Calculate mood score
    mood_match_score = analyze_mood_matches(transcript, moods)

    # Generate personalized improvement suggestions based on actual scores
    improvement_suggestions = generate_targeted_suggestions(
        energy_level_score,
        mood_match_score,
        audio_sequence,
        energy_levels
    )

```


)

Calculate overall score

overall_conductor_score = (energy_level_score + mood_match_score) / 2

result = {

 "transcript": transcript,

 "instructions": instructions,

 "energy_levels": energy_levels,

 "moods": moods,

 "energy_level_score": energy_level_score,

 "mood_match_score": mood_match_score,

 "overall_conductor_score": overall_conductor_score,

 "improvement_suggestions": improvement_suggestions

}

return result

```
xp_system.py
```

```
import math
```

```
from django.db.models import Sum
```

```
from exercises.models import RapidFire, TripleStep, Conductor
```

```
def calculate_user_xp(user):
```

```
    rapidfire_xp =
```

```
    RapidFire.objects.filter(user=user).aggregate(total=Sum('overall_rapidfire_score'))['total'] or  
    0
```

```
    triplestep_xp =
```

```
    TripleStep.objects.filter(user=user).aggregate(total=Sum('overall_triple_step_score'))['total']  
    ] or 0
```

```
    conductor_xp =
```

```
    Conductor.objects.filter(user=user).aggregate(total=Sum('overall_conductor_score'))['total']  
    or 0
```

```
    total_xp = rapidfire_xp + triplestep_xp + conductor_xp
```

```
    return total_xp
```

```
def calculate_user_level(xp, xp_per_level=50):
```

```
    level = math.floor(xp / xp_per_level) + 1
```

```
    return level
```

```

audio_process.py
import os

from pydub import AudioSegment

import whisper

import random

def save_uploaded_file(audio_file_path, destination_path):

    with open(audio_file_path, 'rb') as in_file:

        data = in_file.read()

    with open(destination_path, 'wb') as out_file:

        out_file.write(data)

def convert_wav_to_mp3(wav_file, mp3_file="output.mp3"):

    try:

        audio = AudioSegment.from_wav(wav_file)

        audio.export(mp3_file, format="mp3")

        print(f"Converted {wav_file} to {mp3_file}")

        return mp3_file

    except Exception as e:

        print("Error converting WAV to MP3:", e)

        raise

def transcribe_audio(mp3_file):

    print("Loading Whisper model...")

    model = whisper.load_model("small.en")

    print("Transcribing audio...")

    result = model.transcribe(mp3_file)

    return result

def process_audio_upload(audio_file_path):

    wav_path = "temp_output.wav"

    mp3_path = "temp_output.mp3"

```

try:

```
    print("Converting webm to wav...")
```

```
    audio_segment = AudioSegment.from_file(audio_file_path, format="webm")
```

```
    audio_segment.export(wav_path, format="wav")
```

except Exception as e:

```
    print("Error converting webm to wav:", e)
```

```
    raise
```

```
convert_wav_to_mp3(wav_path, mp3_path)
```

```
transcript = transcribe_audio(mp3_path)
```

```
for file in [wav_path, mp3_path]:
```

```
    if os.path.exists(file):
```

```
        os.remove(file)
```

```
return transcript
```

Emotion Detection in Text.ipynb

EDA

import pandas as pd

import numpy as np

Load Data Viz Pkgs

import seaborn as sns

Load Text Cleaning Pkgs

import neattext.functions as nfx

Load ML Pkgs

Estimators

from sklearn.linear_model import LogisticRegression

from sklearn.naive_bayes import MultinomialNB

Transformers

from sklearn.feature_extraction.text import CountVectorizer

from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

Load Dataset

df = pd.read_csv("emotion_dataset_raw.csv")

df.head()

Value Counts

df['Emotion'].value_counts()

Plot

sns.countplot(x='Emotion', data=df)

Data Cleaning

```
dir(nfx)

# User handles
df['Clean_Text'] = df['Text'].apply(nfx.remove_userhandles)

# Stopwords
df['Clean_Text'] = df['Clean_Text'].apply(nfx.remove_stopwords)

df

# Features & Labels
Xfeatures = df['Clean_Text']

ylabels = df['Emotion']

# Split Data
x_train,x_test,y_train,y_test =
train_test_split(Xfeatures,ylabels,test_size=0.3,random_state=42)

# Build Pipeline
from sklearn.pipeline import Pipeline

# LogisticRegression Pipeline
pipe_lr = Pipeline(steps=[('cv',CountVectorizer()),('lr',LogisticRegression())])

# Train and Fit Data
pipe_lr.fit(x_train,y_train)

pipe_lr

# Check Accuracy
pipe_lr.score(x_test,y_test)

# Make A Prediction
ex1 = "This book was so interesting it made me happy"

pipe_lr.predict([ex1])

# Prediction Prob
pipe_lr.predict_proba([ex1])

pipe_lr.classes_

# Save Model & Pipeline
```

```
import joblib

pipeline_file = open("emotion_classifier_pipe_lr.pkl","wb")
joblib.dump(pipe_lr,pipeline_file)
pipeline_file.close()
```

```
# Load the saved pipeline
```

```
pipeline_file = open("emotion_classifier_pipe_lr.pkl", "rb")
loaded_pipe_lrr = joblib.load(pipeline_file)
pipeline_file.close()
```

```
# Example prediction
```

```
ex2 = "I feel soo sad and lonely today."
predicted_emotion = loaded_pipe_lrr.predict([ex2])
predicted_proba = loaded_pipe_lrr.predict_proba([ex2])
```

```
# Output results
```

```
print(f"Predicted Emotion: {predicted_emotion[0]}")
print(f"Prediction Probabilities: {predicted_proba}")
```